

Average running time for look up on Docker

1 client: 0.030863332748413085
2 clients: 0.08159070014953614
3 clients: 0.11483473777770996
4 clients: 0.15616931915283203
5 clients: 0.18957266807556153

Average running time for trade (or order) on Docker


1 client: 0.13246464729309082
2 clients: 0.26849913597106934
3 clients: 0.38188695907592773
4 clients: 0.49315333366394043
5 clients: 0.5611741542816162

Average running time for look up not on Docker

1 client: 0.024761033058166505
2 clients: 0.04013040065765381
3 clients: 0.061309075355529784
4 clients: 0.08517658710479736
5 clients: 0.08800222873687744

Average running time for trade (or order) not on Docker

1 client: 0.09320902824401855
2 clients: 0.14832735061645508
3 clients: 0.25814008712768555
4 clients: 0.36181092262268066
5 clients: 0.4559962749481201



Answering questions

1. Does the latency of the application change with and without Docker containers? Did virtualization add any overheads?

Latency is smaller without Docker containers. The virtualization adds overheads.

2. How does the latency of the lookup requests compare to trade? Since trade requests involve all these microservices, while lookup requests only involve two microservices, does it impact the observed latency?

Latency of lookup requests is smaller than the one of trade.

It impacts the observed latency.

3. How does the latency change as the number of clients change? Does it change for different types of requests?

The latency increases as the number of clients increases. And it changes for different types of requests.